



Why Threads?

- On single or multiprocessor machines
 - Allow progress when a process is blocked – eg for I/O
 - Convenient programming model
 - Multitasking
- On multiprocessor machines
 - Speedup through parallel processing



What is a Java Thread?

- Conceptually, like a one-shot machine
 - its program is a Runnable:

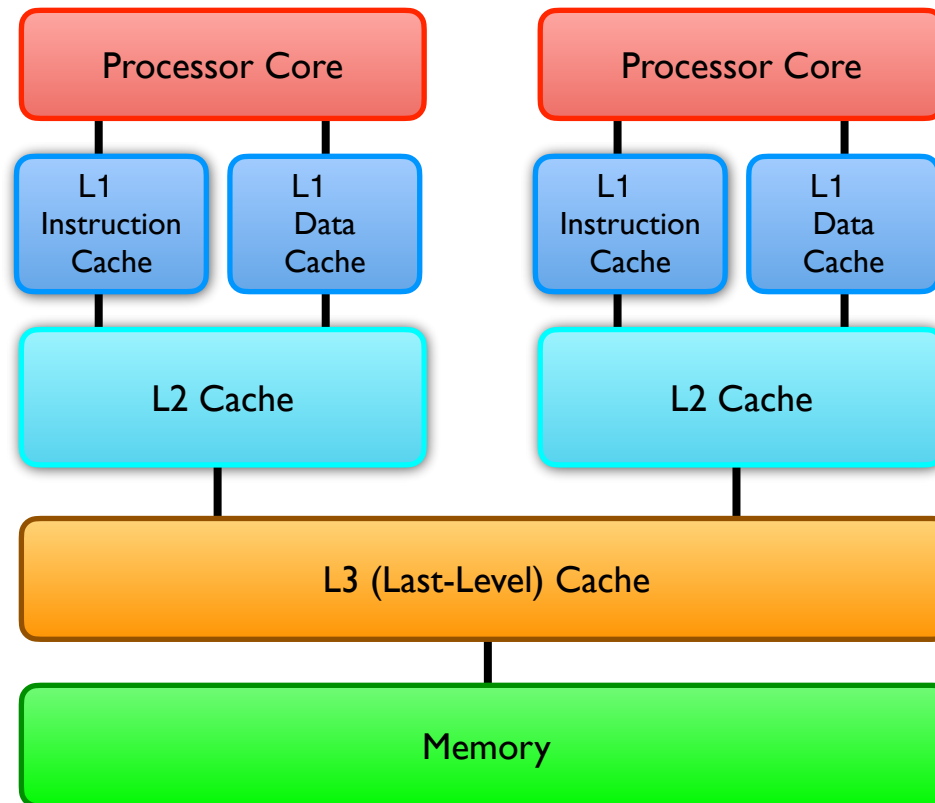
```
public interface Runnable {  
    void run();  
}
```

- A thread dies when its run method terminates
- An application terminates when its last non-daemon thread terminates
- Threads co-operate through shared stored data



Why is Multi-Thread Programming So Hard?

- Modern computer architectures are all about having lots of cores running as fast as possible





java.util.concurrent

- Higher-level utilities that abstract away from the difficult detail of threads:
 - Executor Framework
 - Synchronizers
 - Concurrent Collections
 - Atomic Variables



Possible Executor Implementations

Abstract model of execution context, supports different possibilities:

- In-thread

```
public class InThread implements Executor {  
    public void execute( Runnable task ) {  
        task.run();  
    }  
}
```

- Thread-per-task

```
public class ThreadPerTask implements Executor {  
    public void execute( Runnable task ) {  
        new Thread(task).start();  
    }  
}
```

- Thread-pool...



Executors – Class supplying Executor implementations

- `Executors.newFixedThreadPool()`
 - and `Executors.newSingleThreadExecutor()`
 - fixed set of n threads operating off unbounded queue
 - new threads created to replace threads that crash
- `Executors.newCachedThreadPool()`
 - unlimited size thread pool, new threads created on demand
 - threads unused for 60 seconds terminated and removed
- `Executors.newScheduledThreadPool()`
 - and `Executors.newSingleThreadExecutor()`
 - allows tasks to be scheduled for one-off or repeated execution
- `Executors.newWorkStealingPool()`



Executors – Class supplying Executor implementations

Executors factories actually produce *ExecutorService* implementations

- ExecutorService extends Executor interface
 - methods to manage termination and to track progress of submitted tasks:

```
interface ExecutorService extends Executor {  
    void shutdown();  
    boolean isShutdown();  
    boolean isTerminated();  
    <T> Future<T> submit(Callable<T> task);  
    ...  
}
```



Callable<V> and Future<V>

- Callable<V> is like a Runnable that can return a value and throw exceptions

```
interface Callable<V> {  
    V call() throws Exception  
}
```

- Future<V> represents the progress of an asynchronous computation

```
interface Future<V> {  
    boolean cancel(...);  
    isCancelled();  
    get(...);  
    isDone();  
}
```




Concurrency Exercise 1 – Executor Framework



- Exercise documentation in lab docs folder
- Starting code in `concurrency.DirectoryTreeLISTER`



Synchronizers

- Semaphore
- CountdownLatch
- Also
 - CyclicBarrier
 - Phaser
 - Exchanger
 - FutureTask



- Manages a set of permits
 - used to control the number of activities accessing a resource
 - think of a nightclub bouncer!
- Acquire a permit with `acquire()` or `tryAcquire()`
- Release one or more with `release()`
- Other capabilities: `availablePermits()`, `reducePermits()`
- Basic construct: permits can be lost!



Semaphore throttling task submission

```
public class BoundedExecutor {
    private final Executor exec;
    private final Semaphore semaphore;

    public BoundedExecutor(Executor exec, int bound) {
        this.exec = exec;
        this.semaphore = new Semaphore(bound);
    }

    public void submitTask(final Runnable command) throws InterruptedException {
        semaphore.acquire();
        try {
            exec.execute(new Runnable() {
                public void run() {
                    try {
                        command.run();
                    } finally {
                        semaphore.release();
                    }
                }
            });
        } catch (RejectedExecutionException e) {
            semaphore.release();
        }
    }
}
```



- A *latch* (a one-way gate)
 - Allows one or more threads to wait until a set of events is complete
 - For example, players in a multi-player game can't start until everyone is ready
 - Latch is initialised with a count
 - Threads that call `await()` are blocked until the latch opens
 - Latch opens after `countDown()` has been called enough
 - Subsequent calls to `await()` don't block



CountDownLatch — use in timing test

```
public class TestHarness {
    public long timeTasks(int nThreads, final Runnable task) throws InterruptedException {
        final CountDownLatch startGate = new CountDownLatch(1);
        final CountDownLatch endGate = new CountDownLatch(nThreads);

        for (int i = 0; i < nThreads; i++) {
            Thread t = new Thread() {
                public void run() {
                    try {
                        startGate.await();
                        try {
                            task.run();
                        } finally {
                            endGate.countDown();
                        }
                    } catch (InterruptedException ignored) {}
                }
            };
            t.start();
        }
        long start = System.nanoTime();
        startGate.countDown();
        endGate.await();
        long end = System.nanoTime();
        return end - start;
    }
}
```



Picking a collection:

- If the collection is not shared, non-thread-safe collections are great!
 - `ArrayList`, `LinkedList`, `HashMap`, `TreeMap`
- If the collection is shared but access is not too frequent, use the standard synchronized collections
 - Benefit: ease of use, low memory footprint
- If the collection is shared and frequently accessed by multiple threads, use a concurrent collection
 - `ConcurrentHashMap`, `ConcurrentSkipListMap`, `ConcurrentLinkedQueue`, etc.
 - Benefit: can be used by lots of threads without too much blocking
 - Disadvantage: might use a lot more memory



Synchronized vs Concurrent Collections

- Synchronized collections
 - All methods synchronize on the same lock
 - They hold the lock for the entire operation, even time consuming ones - e.g. `List.contains`
- Concurrent collections can offer dramatic scalability improvements with little risk!
 - At the cost of using more memory
- Most important concurrent collections are `ConcurrentHashMap` and some `BlockingQueue` implementations



ConcurrentHashMap

- Can be used as a drop-in thread-safe replacement for HashMap
- Implementations atomically execute methods that require check-then-act
 - putIfAbsent
 - computeIfAbsent, computeIfPresent
 - merge
 - replace



- Blocking queues are under the hood of every producer-consumer app
 - that is, nearly all concurrent apps
- Blocking queues convey tasks – or data items to be processed – between the processes of a workflow system
 - smoothing out spikes in workload
- Consumers needing a work item should block until one is available
 - that is, the queue is non-empty
- Producers needing to provide a work item should wait until downstream processes are ready for one
 - that is, the queue is non-full



BlockingQueue

```
interface BlockingQueue<E> {  
    void put(E e);  
    E take();  
    boolean offer(E e, long timeout, TimeUnit unit);  
    E poll(long timeout, TimeUnit unit);  
}
```

- Most important implementations are `ArrayBlockingQueue`, `LinkedBlockingQueue`
 - Only `LBQ` can be unbounded
 - Both are `FIFO` structures
- Other implementations include `PriorityBlockingQueue`, `DelayQueue`, `SynchronousQueue`



- Conventional exclusive locking is *pessimistic*
 - Assumes that if you don't guard your valuables, they'll be “rearranged”
 - By someone else!
- With *optimistic* locking you hope everything will be all right
 - If it isn't, your operation failed to finish without interference
 - You can just try it again
- Optimistic locking relies on hardware support – often a compare-and-swap (CAS) instruction
- Can be much more efficient than pessimistic locking, especially when contention isn't high



How CAS works

```
public class SimulatedCAS {  
    private int value;  
    public synchronized int get() { return value; }  
    public synchronized int compareAndSwap(int expectedValue,  
                                           int newValue) {  
        int oldValue = value;  
        if (oldValue == expectedValue)  
            value = newValue;  
        return oldValue;  
    }  
}
```

```
SimulatedCAS cas = ...  
valueToWorkOn = cas.get();  
newValue = doLongRunningOperation(valueToWorkOn);  
cas.compareAndSwap(valueToWorkOn, newValue)
```



- So-called because its operations are atomic
- Pseudo-code for addAndGet:

```
public final int addAndGet(int delta) {  
    for (;;) {  
        int current = get();  
        int next = current + delta;  
        if (compareAndSwap(current, next) == current)  
            return next;  
    }  
}
```



Atomics as “Better Volatiles”

- Atomic variables are thread-safe without synchronisation
 - Values stored internally as volatile fields
 - Same visibility semantics
 - Little reason now to use volatile directly



Types of Atomic classes

- The following types have atomics built in
 - AtomicBoolean
 - AtomicInteger
 - Use for int, short, byte and float (use `Float.floatToIntBits(float)`)
- AtomicLong
 - Use for long and double
- AtomicReference
- There are also atomic array classes
 - Necessary as you can never make values of an array volatile!
 - AtomicIntegerArray, AtomicLongArray, AtomicReferenceArray



- Fast counter when we have high contention
 - Stripes the values into different cells to reduce contention

```
public class BankAccount {  
    private final LongAdder balance = new LongAdder();  
    public BankAccount(long balance) {  
        this.balance.add(balance);  
    }  
    public void deposit(long amount) {  
        balance.add(amount);  
    }  
    public void withdraw(long amount) {  
        deposit(-amount);  
    }  
    public long getBalance() {  
        return balance.longValue();  
    }  
}
```



Concurrency Exercise 2



- Exercise documentation:
- Starting code: